

# Kernel-Hack-Drill: Environment For Developing Linux Kernel Exploits


Alexander Popov

■ **positive technologies**



April 10, 2025

# Teaser



# Who Am I

- Alexander Popov
- Linux kernel developer since 2012
- Maintainer of some free software projects
- Principal Security Researcher and Head of  
    Open Source Program Office at  **positive technologies**
- Conference speaker:

Zer0Con, OffensiveCon, H2HC, Nullcon Goa, Linux Security Summit, Still Hacking Anyway, HITB,  
Positive Hack Days, ZeroNights, HighLoad++, Open Source Summit, OS Day, Linux Plumbers...

[a13xp0p0v.github.io/conference\\_talks](https://a13xp0p0v.github.io/conference_talks)



# Agenda

- ① The bug collision story
- ② About **CVE-2024-50264**
- ③ A new approach to exploiting it
- ④ How **kernel-hack-drill** helped to achieve this



# How It Began

- I first found and exploited a bug in **AF\_VSOCK** in 2021:  
Four Bytes of Power: Exploiting CVE-2021-26708 in the Linux kernel  
[a13xp0p0v.github.io/2021/02/09/CVE-2021-26708.html](https://a13xp0p0v.github.io/2021/02/09/CVE-2021-26708.html)
- In spring 2024, I was fuzzing the kernel with a customized syzkaller
- I found another bug in **AF\_VSOCK** in April 2024
- I minimized the reproducer, disabled KASAN and  
got instant null-ptr-deref in a kernel worker
- Postponed this bug

# Bug Collision

- I decided to look at this bug again in autumn 2024
- Results were promising but then...

# Bug Collision

- I decided to look at this bug again in autumn 2024
- Results were promising but then...
- Got bug collision with Hyunwoo Kim ([@v4bel](#)) and Wongi Lee ([@qwerty](#))
- They disclosed this bug as **CVE-2024-50264** and used it at kernelCTF
- Their patch turned my PoC into null-ptr-deref

```
Diffstat (limited to 'net/vmw_vsock')
-rw-r-- net/vmw_vsock/virtio_transport_common.c 1
1 files changed, 1 insertions, 0 deletions

diff --git a/net/vmw_vsock/virtio_transport_common.c b/net/vmw_vsock/virtio_transport_common.c
index ccbd2bc0d2109a..fc5666c8298f7b 100644
--- a/net/vmw_vsock/virtio_transport_common.c
+++ b/net/vmw_vsock/virtio_transport_common.c
@@ -1109,6 +1109,7 @@ void virtio_transport_destruct(struct vsock_sock *vsk)
    struct virtio_vsock_sock *vvs = vsk->trans;
    kfree(vvs);
+   vsk->trans = NULL;
}
EXPORT_SYMBOL_GPL(virtio_transport_destruct);
```

# Continue Anyway

- The exploit strategy by `@v4bel` and `@qwerty` looked very complicated  
[github.com/google/security-research/pull/145/files](https://github.com/google/security-research/pull/145/files)
- I had some different ideas and decided to continue my research anyway
- I chose **Ubuntu Server 24.04** with a fresh OEM/HWE kernel (**v6.11**) as the target



Viktor Vasnetsov: The Knight at the Crossroads (1878)

- The bug was introduced in August 2016 (commit 06a8fc78367d, Linux v4.8)
- Race condition in `AF_VSOCK` sockets between `connect()` and a POSIX signal
- `CONFIG_USER_NS` is not required
- UAF on `virtio_vsock_sock` object (`kmalloc-96`)
- Memory corruption: UAF write in a kernel worker
- It has a lot of nasty limitations for the exploitation
  - **The worst bug** for the exploitation that I've ever seen

# Reproducing CVE-2024-50264: Immortal Signal Handler

- @v4bel & @qwerty used **SIGKILL**
- My fuzzer found another approach, which amazed me

```
struct sigevent sev = {};
timer_t race_timer = 0;
sev.sigev_notify = SIGEV_SIGNAL; /* Notification type */
sev.sigev_signo = 33; /* Secret NPTL Signal (see nptl(7)) */
ret = timer_create(CLOCK_MONOTONIC, &sev, &race_timer);
```



- Native POSIX Threads Library makes internal use of signal 33
- Syscall wrappers and glibc functions **hide this signal** from applications
- So I can use `timer_settime()` for `race_timer`
  - It gives **control of timing**: at which moment signal should interrupt `connect()`
  - It is invisible for the exploit process and **doesn't kill it**

# CVE-2024-50264: Code Performing UAF Write

- This function is called **in kworker** after `virtio_vsock_sock` is freed

```
static bool virtio_transport_space_update(struct sock *sk,
                                         struct sk_buff *skb)
{
    struct virtio_vsock_hdr *hdr = virtio_vsock_hdr(skb);
    struct vsock_sock *vsk = vsock_sk(sk);
    struct virtio_vsock_sock *vvs = vsk->trans;           /* ptr to freed object */
    bool space_available;

    if (!vvs)
        return true;


    spin_lock_bh(&vvs->tx_lock); /* proceed if 4 bytes are zero (UAF write non-zero to lock) */
    vvs->peer_buf_alloc = le32_to_cpu(hdr->buf_alloc); /* UAF write 4 bytes */
    vvs->peer_fwd_cnt = le32_to_cpu(hdr->fwd_cnt);   /* UAF write 4 bytes */
    space_available = virtio_transport_has_space(vsk); /* UAF read, not interesting */
    spin_unlock_bh(&vvs->tx_lock);                   /* UAF write, restore 4 zero bytes */
    return space_available;
}
```



- There is no pointer dereference in freed object

# CVE-2024-50264: UAF Write

**struct virtio\_vsock\_sock**



# UAF Write: Data Control

- About `virtio_vsock_sock.peer_buf_alloc` value control from userspace:

```
/* Increase the range for the value that we want to write during UAF: */
uaf_val_limit = 0x1lu; /* can't be zero */
setsockopt(vsck1, PF_VSOCK, SO_VM_SOCKETS_BUFFER_MIN_SIZE,
           &uaf_val_limit, sizeof(uaf_val_limit));
uaf_val_limit = 0xffffffffflu;
setsockopt(vsck1, PF_VSOCK, SO_VM_SOCKETS_BUFFER_MAX_SIZE,
           &uaf_val_limit, sizeof(uaf_val_limit));

/* Set the 4-byte value that we want to write during UAF: */
setsockopt(vsck1, PF_VSOCK, SO_VM_SOCKETS_BUFFER_SIZE,
           &uaf_val, sizeof(uaf_val));
```

- About `virtio_vsock_sock.peer_fwd_cnt` value control from userspace:

- It represents the number of bytes pushed through `vsock` using `sendmsg()`/`recvmsg()`
- Zero by default (4 zero bytes)

# Not So Fast: CVE-2024-50264 Limitations

- ➊ Vulnerable `virtio_vsock_sock` client object is allocated together with the server one

- It's bad for cross-cache attack



<https://www.youtube.com/watch?v=hbKEdmPPxy4>

# Not So Fast: CVE-2024-50264 Limitations

- ➊ Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
  - It's bad for cross-cache attack
- ➋ Reproducing this race condition is very unstable



<https://www.youtube.com/watch?v=hbKEdmPPxy4>

# Not So Fast: CVE-2024-50264 Limitations

- ➊ Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
  - It's bad for cross-cache attack
- ➋ Reproducing this race condition is very unstable
- ➌ UAF write happens in kworker within few `µs` after `kfree()`



<https://www.youtube.com/watch?v=hbKEdmPPxy4>

# Not So Fast: CVE-2024-50264 Limitations

- ➊ Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
  - It's bad for cross-cache attack
- ➋ Reproducing this race condition is very unstable
- ➌ UAF write happens in kworker within few `µs` after `kfree()`
- ➍ Null-ptr-deref happens in kworker right after UAF write



<https://www.youtube.com/watch?v=hbKEdmPPxy4>

# Not So Fast: CVE-2024-50264 Limitations

- ➊ Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
  - It's bad for cross-cache attack
- ➋ Reproducing this race condition is very unstable
- ➌ UAF write happens in kworker within few `µs` after `kfree()`
- ➍ Null-ptr-deref happens in kworker right after UAF write
- ➎ If this kernel oops is avoided, another null-ptr-deref happens in kworker after `VSOCK_CLOSE_TIMEOUT` (8 sec)



<https://www.youtube.com/watch?v=hbKEdmPPxy4>

# Not So Fast: CVE-2024-50264 Limitations

- ➊ Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
  - It's bad for cross-cache attack
- ➋ Reproducing this race condition is very unstable
- ➌ UAF write happens in kworker within few `µs` after `kfree()`
- ➍ Null-ptr-deref happens in kworker right after UAF write
- ➎ If this kernel oops is avoided, another null-ptr-deref happens in kworker after `VSOCK_CLOSE_TIMEOUT` (8 sec)
- ➏ Kworker hangs if `virtio_vsock_sock.tx_lock` is not zero



<https://www.youtube.com/watch?v=hbKEdmPPxy4>

# Not So Fast: CVE-2024-50264 Limitations



## Challenge

Now you can see why this was the worst bug  
for exploitation I had ever seen

# Exploit Strategy of @v4bel & @qwerty


- ① Large-scale BPF JIT Spray populating a significant portion of the physical memory



# Exploit Strategy of @v4bel & @qwerty

- ① Large-scale BPF JIT Spray populating a significant portion of the physical memory
- ② SLUBStick technique [github.com/IAIK/SLUBStick](https://github.com/IAIK/SLUBStick)

- Using timing side channel to determine number of objects in the active slab
- Allocating the `virtio_vsock_sock` client and server objects in different slabs
- It's possible by making them the `last` and `first` objects in slabs



WOW LOOKS COMPLICATED



CAN I DO IT DIFFERENTLY?

# Exploit Strategy of @v4bel & @qwerty

① Large-scale BPF JIT Spray populating a significant portion of the physical memory

② SLUBStick technique [github.com/IAIK/SLUBStick](https://github.com/IAIK/SLUBStick)

- Using timing side channel to determine number of objects in the active slab
- Allocating the `virtio_vsock_sock` client and server objects in different slabs
- It's possible by making them the `last` and `first` objects in slabs

③ Dirty Pagetable technique [yanglingxi1993.github.io/dirty\\_pagetable/dirty\\_pagetable.html](https://yanglingxi1993.github.io/dirty_pagetable/dirty_pagetable.html)

- Cross-allocator attack reclaiming slab with UAF object for Page Table Entry
- UAF write to PTE to make it possibly point a BPF JIT region



# Exploit Strategy of @v4bel & @qwerty

① Large-scale BPF JIT Spray populating a significant portion of the physical memory

② SLUBStick technique [github.com/IAIK/SLUBStick](https://github.com/IAIK/SLUBStick)

- Using timing side channel to determine number of objects in the active slab
- Allocating the `virtio_vsock_sock` client and server objects in different slabs
- It's possible by making them the `last` and `first` objects in slabs

③ Dirty Pagetable technique [yanglingxi1993.github.io/dirty\\_pagetable/dirty\\_pagetable.html](https://yanglingxi1993.github.io/dirty_pagetable/dirty_pagetable.html)

- Cross-allocator attack reclaiming slab with UAF object for Page Table Entry
- UAF write to PTE to make it possibly point a BPF JIT region

④ Inserting the privilege escalation payload into BPF code



# Exploit Strategy of @v4bel & @qwerty

① Large-scale BPF JIT Spray populating a significant portion of the physical memory

② SLUBStick technique [github.com/IAIK/SLUBStick](https://github.com/IAIK/SLUBStick)

- Using timing side channel to determine number of objects in the active slab
- Allocating the `virtio_vsock_sock` client and server objects in different slabs
- It's possible by making them the `last` and `first` objects in slabs



③ Dirty Pagetable technique [yanglingxi1993.github.io/dirty\\_pagetable/dirty\\_pagetable.html](https://yanglingxi1993.github.io/dirty_pagetable/dirty_pagetable.html)

- Cross-allocator attack reclaiming slab with UAF object for Page Table Entry
- UAF write to PTE to make it possibly point a BPF JIT region



④ Inserting the privilege escalation payload into BPF code

⑤ Socket communication to trigger the privesc payload

# My First Ideas on Exploit Strategy


- Try UAF write to some kernel object
- Should I search kernel objects in `kmalloc-96`?
- **No!** Ubuntu Server 24.04 has:
  - `CONFIG_RANDOM_KMALLOC_CACHES=y`
  - `CONFIG_SLAB_BUCKETS=y`
  - `CONFIG_SLUB_CPU_PARTIAL=y`
- I will try cross-cache attack

# Possible Target for UAF Write: struct cred

struct virtio_vssock_sock (80 bytes, 2 items)		struct cred (184 bytes)	
0		0	
8	spinlock_t tx_lock==0	8	kuid_t uid != 0
16		16	gid_t sgid
24	u32 peer_buf_alloc	24	kuid_t euid
32		32	
40		40	
48		48	
56		56	
64		64	
72		72	
80	8-byte hole	80	
88	8-byte hole	88	
96		96	
104	spinlock_t tx_lock==0	104	struct key *process_keyring
112		112	struct key *thread_keyring
120	u32 peer_buf_alloc	120	struct key *request_key_auth
128		128	
136		136	
144		144	
152		152	
160		160	
168		168	
176	8-byte hole	176	

total size: 184 bytes

# Target for UAF Write: struct cred (No Way)



# Target for UAF Write: struct msg\_msg

- Why? Because I like it
- I first used it as a UAF target object in 2021

[a13xp0p0v.github.io/2021/02/09/CVE-2021-26708.html](https://a13xp0p0v.github.io/2021/02/09/CVE-2021-26708.html)

- It was a novel approach back then
- I decided to create something new again



# virtio\_vsock\_sock vs msg\_msg

**struct virtio\_vsock\_sock (80 bytes)**

0	
8	spinlock_t tx_lock==0
16	
24	u32 peer_buf_alloc
32	
40	
48	
56	
64	
72	
80	8-byte hole
88	8-byte hole

kernel hang  
on spinlock



**struct msg\_msg (96 bytes)**

0	
8	struct list_head *prev != 0
16	long int m_type
24	size_t m_ts
32	
40	
48	
56	
64	
72	
80	
88	



# Bypassing the Unwanted msg\_msg.m\_list Corruption

- `msg_msg.m_list.prev` would be interpreted as non-null `tx_lock`
- `virtio_transport_space_update()` would hang in `spin_lock_bh()`
- Need to initialize `msg_msg.m_list.prev` **after** the UAF write
- Can we postpone placing `msg_msg` in the message queue?
- Yes!

# Spray msg\_msg Allowing m\_list Corruption (Novel Technique?)

- ➊ Fill the message queue almost completely before sending the target `msg_msg`

- The message queue size is `MSGMNB` (16384 bytes)
- Send 2 clogging messages of 8191 bytes each
- 2 bytes left in the queue, don't call `msgrecv()`



<https://www.youtube.com/watch?v=0XVCz6nekJc>

# Spray msg\_msg Allowing m\_list Corruption (Novel Technique?)

## ① Fill the message queue almost completely before sending the target msg\_msg

- The message queue size is MSGMNB (16384 bytes)
- Send 2 clogging messages of 8191 bytes each
- 2 bytes left in the queue, don't call msgrcv()

## ② Spray target msg\_msg objects

- Call the msgsnd() syscall in separate pthreads
- Kernel allocates target msg\_msg and msgsnd() blocks



<https://www.youtube.com/watch?v=0XVCz6nekJc>

# Spray msg\_msg Allowing m\_list Corruption (Novel Technique?)

① Fill the message queue almost completely before sending the target `msg_msg`

- The message queue size is `MSGMNB` (16384 bytes)
- Send 2 clogging messages of 8191 bytes each
- 2 bytes left in the queue, don't call `msgrecv()`

② Spray target `msg_msg` objects

- Call the `msgsnd()` syscall in separate pthreads
- Kernel allocates target `msg_msg` and `msgsnd()` blocks

③ Perform UAF write, corrupt `msg_msg.m_list` as you want



<https://www.youtube.com/watch?v=0XVCz6nekJc>

# Spray msg\_msg Allowing m\_list Corruption (Novel Technique?)

① Fill the message queue almost completely before sending the target `msg_msg`

- The message queue size is `MSGMNB` (16384 bytes)
- Send 2 clogging messages of 8191 bytes each
- 2 bytes left in the queue, don't call `msgrecv()`

② Spray target `msg_msg` objects

- Call the `msgsnd()` syscall in separate pthreads
- Kernel allocates target `msg_msg` and `msgsnd()` blocks

③ Perform UAF write, corrupt `msg_msg.m_list` as you want

④ Perform `msgrecv()` for clogging messages

- Now the kernel can add sprayed `msg_msg` to the queue
- The **kernel fixes** the **corrupted** `msg_msg.m_list` pointers!



<https://www.youtube.com/watch?v=0XVCz6nekJc>

# virtio\_vsock\_sock vs msg\_msg

**struct virtio\_vsock\_sock (80 bytes)**

0	
8	spinlock_t tx_lock==0
16	
24	u32 peer_buf_alloc
32	
40	
48	
56	
64	
72	
80	8-byte hole
88	8-byte hole

**struct msg\_msg (96 bytes)**

0	
8	struct list_head *prev == 0
16	long int m_type
24	size_t m_ts
32	
40	
48	
56	
64	
72	
80	
88	



# Nice Trick, What's Next?

- ➊ I managed to overwrite `msg_msg.m_ts` and make kernel fix up `msg_msg.m_list`
  - This technique is also useful for **blind overwriting of msg\_msg**
  - No kernel infoleak is needed — the kernel will restore the corrupted pointers

# Nice Trick, What's Next?

- ➊ I managed to overwrite `msg_msg.m_ts` and make kernel fix up `msg_msg.m_list`
  - This technique is also useful for **blind overwriting of msg\_msg**
  - No kernel infoleak is needed — the kernel will restore the corrupted pointers
- ➋ To use this trick, I needed to perform cross-cache attack
  - `virtio_vsock_sock` lives in one of 16 `kmalloc-rnd-?-96` slab caches (`CONFIG_RANDOM_KMALLOC_CACHES`)
  - `msg_msg` lives in `msg_msg-96` slab cache (`CONFIG_SLAB_BUCKETS`)

# Nice Trick, What's Next?

- ➊ I managed to overwrite `msg_msg.m_ts` and make kernel fix up `msg_msg.m_list`
  - This technique is also useful for **blind overwriting of msg\_msg**
  - No kernel infoleak is needed — the kernel will restore the corrupted pointers
- ➋ To use this trick, I needed to perform cross-cache attack
  - `virtio_vsock_sock` lives in one of 16 `kmalloc-rnd-?-96` slab caches (`CONFIG_RANDOM_KMALLOC_CACHES`)
  - `msg_msg` lives in `msg_msg-96` slab cache (`CONFIG_SLAB_BUCKETS`)
- ➌ **Problems:**
  - I needed to learn how cross-cache attacks work on the latest Ubuntu kernel
  - Testing exploit primitives together with this crazy race condition was **painful**

# Solution That Makes Researcher's Life Easier



Unstable race condition creating problems?

Use a testing ground for developing  
the exploit primitives!

# Kernel Hack Drill

- Open-source project: [github.com/a13xp0p0v/kernel-hack-drill](https://github.com/a13xp0p0v/kernel-hack-drill)
- Provides test environment for developing the Linux kernel exploit primitives you need
- Includes a good step-by-step setup guide in the README (kudos to the contributors!)
- A bit similar to [github.com/hacktivesec/KRWX](https://github.com/hacktivesec/KRWX), but
  - Much simpler
  - Contains interesting PoC exploits



<https://www.pngall.com/wp-content/uploads/4/Drill-Machine-PNG-Free-Download.png>

# Kernel Hack Drill Contents: Kernel Module

## ① `drill_mod.c`

- A small Linux kernel module
- Provides `/proc/drill_act` file as a simple interface to userspace
- Contains nice vulnerabilities that you control

## ② `drill.h`

- Header file describing the `drill_mod.ko` interface

## ③ `drill_test.c`

- Userspace test for `drill_mod.ko`
- It also passes if `CONFIG_KASAN=y`

```
#define DRILL_N 10240
#define DRILL_ITEM_SIZE 95

struct drill_item_t {
    unsigned long foobar;
    void (*callback)(void);
    char data[]; /* C99 flexible array */
};

enum drill_act_t {
    DRILL_ACT_NONE = 0,
    DRILL_ACT_ALLOC = 1,
    DRILL_ACT_CALLBACK = 2,
    DRILL_ACT_SAVE_VAL = 3,
    DRILL_ACT_FREE = 4,
    DRILL_ACT_RESET = 5
};
```

# Kernel Hack Drill Contents: PoC Exploits

## ① `drill_uaf_callback.c`

- UAF exploit invoking a callback in the freed `drill_item_t` struct
- Performs control flow hijack and gains LPE



<https://www.printables.com/model/78077-drill-guide>

# Kernel Hack Drill Contents: PoC Exploits

## ① `drill_uaf_callback.c`

- UAF exploit invoking a callback in the freed `drill_item_t` struct
- Performs control flow hijack and gains LPE

## ② `drill_uaf_write_msg_msg.c`

- UAF exploit writing data to the freed `drill_item_t` struct
- Performs a cross-cache attack, overwrites `msg_msg.m_ts`
- Enables out-of-bounds read of the kernel memory



<https://www.printables.com/model/78077-drill-guide>

# Kernel Hack Drill Contents: PoC Exploits

## ① `drill_uaf_callback.c`

- UAF exploit invoking a callback in the freed `drill_item_t` struct
- Performs control flow hijack and gains LPE

## ② `drill_uaf_write_msg_msg.c`

- UAF exploit writing data to the freed `drill_item_t` struct
- Performs a cross-cache attack, overwrites `msg_msg.m_ts`
- Enables out-of-bounds read of the kernel memory

## ③ `drill_uaf_write_pipe_buffer.c`

- UAF exploit writing data to the freed `drill_item_t` struct
- Performs cross-cache attack, overwrites `pipe_buffer.flags`
- Implements the Dirty Pipe attack and gains LPE



<https://www.printables.com/model/78077-drill-guide>

# Kernel Hack Drill Contents: PoC Exploits

## ① `drill_uaf_callback.c`

- UAF exploit invoking a callback in the freed `drill_item_t` struct
- Performs control flow hijack and gains LPE

## ② `drill_uaf_write_msg_msg.c`

- UAF exploit writing data to the freed `drill_item_t` struct
- Performs a cross-cache attack, overwrites `msg_msg.m_ts`
- Enables out-of-bounds read of the kernel memory

## ③ `drill_uaf_write_pipe_buffer.c`

- UAF exploit writing data to the freed `drill_item_t` struct
- Performs cross-cache attack, overwrites `pipe_buffer.flags`
- Implements the Dirty Pipe attack and gains LPE

## ④ More PoC exploits will come soon!



<https://www.printables.com/model/78077-drill-guide>

# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects
- ⑤ Obtain dangling reference to the vulnerable object for UAF



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects
- ⑤ Obtain dangling reference to the vulnerable object for UAF
- ⑥ Create a new active slab: allocate `objs_per_slab` objects



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects
- ⑤ Obtain dangling reference to the vulnerable object for UAF
- ⑥ Create a new active slab: allocate `objs_per_slab` objects
- ⑦ Free the slab with UAF object: free `(objs_per_slab * 2 - 1)` objects before the last one



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects
- ⑤ Obtain dangling reference to the vulnerable object for UAF
- ⑥ Create a new active slab: allocate `objs_per_slab` objects
- ⑦ Free the slab with UAF object: free `(objs_per_slab * 2 - 1)` objects before the last one
- ⑧ Clean up the partial list: free one of each `objs_per_slab` objects in the reserved slabs



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects
- ⑤ Obtain dangling reference to the vulnerable object for UAF
- ⑥ Create a new active slab: allocate `objs_per_slab` objects
- ⑦ Free the slab with UAF object: free `(objs_per_slab * 2 - 1)` objects before the last one
- ⑧ Clean up the partial list: free one of each `objs_per_slab` objects in the reserved slabs
- ⑨ Reclaim the page with UAF object: spray target objects



# Cross-Cache Attack in Kernel Hack Drill

Standard cross-cache procedure, see the code: [kernel-hack-drill/drill\\_uaf\\_write\\_msg\\_msg.c](#)

- ① Collect the needed info in `/sys/kernel/slab`: `cpu_partial=120, objs_per_slab=42`
- ② Create a new active slab: allocate `objs_per_slab` objects
- ③ Allocate `(objs_per_slab * cpu_partial)` objects for the partial list
- ④ Create the vulnerable slab: allocate `objs_per_slab` objects
- ⑤ Obtain dangling reference to the vulnerable object for UAF
- ⑥ Create a new active slab: allocate `objs_per_slab` objects
- ⑦ Free the slab with UAF object: free `(objs_per_slab * 2 - 1)` objects before the last one
- ⑧ Clean up the partial list: free one of each `objs_per_slab` objects in the reserved slabs
- ⑨ Reclaim the page with UAF object: spray target objects
- ⑩ Exploit UAF



# Debugging Cross-Cache Attack: Kernel Patch

```
diff --git a/ipc/msgutil.c b/ipc/msgutil.c
@@ -64,6 +64,7 @@ static struct msg_msg *alloc_msg(size_t len)
     msg = kmem_buckets_alloc(msg_buckets, sizeof(*msg) + alen, GFP_KERNEL);
     if (msg == NULL)
         return NULL;
+    printk("msg_msg 0x%lx\n", (unsigned long)msg);

     msg->next = NULL;
     msg->security = NULL;

diff --git a/mm/slub.c b/mm/slub.c
@@ -3140,6 +3140,7 @@ static void __put_partials(struct kmem_cache *s, struct slab *partial_slab)
     while (slab_to_discard) {
         slab = slab_to_discard;
         slab_to_discard = slab_to_discard->next;
+        printk("__put_partials: cache 0x%lx slab 0x%lx\n", (unsigned long)s, (unsigned long)slab);

         stat(s, DEACTIVATE_EMPTY);
         discard_slab(s, slab);
```

- `__put_partials()` calls `discard_slab()`, which moves the slab to the page allocator

# Debugging Cross-Cache Attack: Console Output and GDB

- Legend: kernel log, stdout, GDB session (with bata24/gef)

```
[ 49.755740] drill: kmalloc'ed item 5081 (0xfffff8880068878a0, size 95)
[+] current_n: 5082 (next for allocating)
4) obtain dangling reference from use-after-free bug
[+] uaf_n: 5081

gef> slab-contains 0xfffff8880068878a0
[+] Wait for memory scan
slab: 0xfffffea00001a21c0
kmem_cache: 0xfffff88800384e800
base: 0xfffff888006887000
name: kmalloc-rnd-14-96 size: 0x60 num_pages: 0x1

[ 51.371255] __put_partials: cache 0xfffff88800384e800 slab 0xfffffea00001a21c0
[ 51.463570] msg_msg 0xfffff8880068878a0
```

- The `drill_item_t` object `0xfffff8880068878a0` in slab `0xfffffea00001a21c0` is reallocated as `msg_msg`

# In My Humble Opinion



**RECENT  
SLAB HARDENING  
FEATURES**

**KERNEL FEATURES  
THAT MAKE  
CROSS-CACHE ATTACKS  
COMPLETELY STABLE**

imgflip.com

# Cross-Cache Attack: Adoption to AF\_VSOCK Exploit

- The vulnerable `virtio_vsock_sock` client object is allocated together with the server one
- It is harmful for the attack (**Limitation #1**):
  - Not closing server vsock prevents complete freeing of UAF slab
  - Closing server vsock breaks UAF
- How can we cope with it?
  - @v4bel and @qwerty used the **SLUBStick** technique

# Cross-Cache Attack: Adoption to AF\_VSOCK Exploit

- The vulnerable `virtio_vsock_sock` client object is allocated together with the server one
- It is harmful for the attack (**Limitation #1**):
  - Not closing server `vsock` prevents complete freeing of UAF slab
  - Closing server `vsock` breaks UAF
- How can we cope with it?
  - @v4bel and @qwerty used the **SLUBStick** technique
  - **My idea:** what if we hit `connect()` with a signal **very early**?

# Race Conditions Are Awful/Awesome

I used one more race condition to exploit the main race condition

- ① Hit `vsock connect()` with the "immortal" signal 33 after 10000 ns
- ② Check whether the race condition succeeded:
  - The `connect()` syscall should return "Interrupted system call"
  - Connecting to server `vsock` from another test client `vsock` should succeed
- ③ If that is true, only a single vulnerable `vsock` was created
- ④ ~~Limitation #1~~ (paired object creation) is bypassed
- ⑤ Cool, the cross-cache attack for `vsock` is unlocked!



# AF\_VSOCK Exploit Speedrun

- This smart testing of `signal` vs `connect()` state also made the exploit **much faster**
  - The UAF write can now be triggered **once per second** instead of ~~once per several minutes~~
  - **Limitation #2** (unstable race condition) is mitigated
  - **Limitation #5** (kworker oops in 8 sec) is bypassed
- To counter **Limitation #4** (kworker oops just after UAF), I used one more race condition
  - Idea by @v4bel and @qwerty
  - Call `listen()` for vulnerable `vsock` just after `connect()` provoking UAF
  - If we are lucky, `listen()` executes before UAF-kworker and prevents null-ptr-deref
  - This is the **main source of instability** of the whole exploit 😕

# Not So Fast: CVE-2024-50264 Limitations

- ① Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
- ② Reproducing this race condition is very unstable
- ③ UAF write happens in kworker within few `µs` after `kfree()`
- ④ Null-ptr-deref happens in kworker right after UAF write
- ⑤ If this kernel oops is avoided, another null-ptr-deref happens  
in kworker after `VSOCK_CLOSE_TIMEOUT` (8 sec)
- ⑥ Kworker hangs if `virtio_vsock_sock.tx_lock` is not zero



## Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow

# Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow
- To deal with **Limitation #3**, I also used a well-known technique by **Jann Horn**  
[googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html](http://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html)
- Hit kworker with a timer interrupt that has **a lot of** `epoll` watches registered for `timerfd`

## Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow
- To deal with **Limitation #3**, I also used a well-known technique by **Jann Horn**  
[googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html](http://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html)
- Hit kworker with a timer interrupt that has **a lot of** `epoll` watches registered for `timerfd`
  - ➊ Call `timerfd_create(CLOCK_MONOTONIC, 0)`

# Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow
- To deal with **Limitation #3**, I also used a well-known technique by **Jann Horn**  
[googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html](http://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html)
- Hit kworker with a timer interrupt that has **a lot of** `epoll` watches registered for `timerfd`
  - ➊ Call `timerfd_create(CLOCK_MONOTONIC, 0)`
  - ➋ Create 8 forks, call `dup()` 100 times for `timertfd` in each fork

# Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow
- To deal with **Limitation #3**, I also used a well-known technique by **Jann Horn**  
[googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html](http://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html)
- Hit kworker with a timer interrupt that has **a lot of** `epoll` watches registered for `timerfd`
  - ➊ Call `timerfd_create(CLOCK_MONOTONIC, 0)`
  - ➋ Create 8 forks, call `dup()` 100 times for `timertfd` in each fork
  - ➌ Call `epoll_create()` 500 times in each fork, register `epoll_fd` for each duplicated `fd`

# Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow
- To deal with **Limitation #3**, I also used a well-known technique by **Jann Horn**  
[googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html](http://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html)
- Hit kworker with a timer interrupt that has **a lot of** `epoll` watches registered for `timerfd`
  - ➊ Call `timerfd_create(CLOCK_MONOTONIC, 0)`
  - ➋ Create 8 forks, call `dup()` 100 times for `timertfd` in each fork
  - ➌ Call `epoll_create()` 500 times in each fork, register `epoll_fd` for each duplicated `fd`
  - ➍ Don't exceed `/proc/sys/fs/epoll/max_user_watches` ( $8 \times 100 \times 500 < 446976$ )

# Not So Fast: Cross-Cache Attack is Too Late

- UAF write in kworker happens within few `us` after `kfree(virtio_vsock_sock)`
- The cross-cache attack is too slow
- To deal with **Limitation #3**, I also used a well-known technique by **Jann Horn**  
[googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html](http://googleprojectzero.blogspot.com/2022/03/racing-against-clock-hitting-tiny.html)
- Hit kworker with a timer interrupt that has **a lot of** `epoll` watches registered for `timerfd`
  - ➊ Call `timerfd_create(CLOCK_MONOTONIC, 0)`
  - ➋ Create 8 forks, call `dup()` 100 times for `timertfd` in each fork
  - ➌ Call `epoll_create()` 500 times in each fork, register `epoll_fd` for each duplicated `fd`
  - ➍ Don't exceed `/proc/sys/fs/epoll/max_user_watches` ( $8 \times 100 \times 500 < 446976$ )
  - ➎ Shoot into kworker setting the proper timeout:  
`timerfd_settime(timerfd, TFD_TIMER_CANCEL_ON_SET, &retard_tmo, NULL)`

# Achieved msg\_msg Out-Of-Bounds Read

- vsock UAF changes the msg\_msg data size from 48 bytes to 8192 (MSGMAX)
- Cool, msgrcv() performs out-of-bounds read of kernel memory
- What does infoleak provide?
  - A kernel address 0xffffffff8233cfa0
  - GDB shows that it is pointer to socket\_file\_ops()
  - Which kernel object stores it? It's struct file!
  - It contains f\_cred pointer, which also leaked
- This infoleak works with high probability



artifact from Hermitage

WHAT'S INSIDE?

## What's Next?



The most interesting / difficult part of the research

Then I needed arbitrary address writing  
for privilege escalation.

I wanted to implement data-only attack  
without control flow hijacking.

## How About Dirty Page Table Attack?

- Good description: [yanglingxi1993.github.io/dirty\\_pagetable/dirty\\_pagetable.html](https://yanglingxi1993.github.io/dirty_pagetable/dirty_pagetable.html)
- Attacking page tables requires **knowing the physical address** of kernel text/heap

## How About Dirty Page Table Attack?

- Good description: [yanglingxi1993.github.io/dirty\\_pagetable/dirty\\_pagetable.html](https://yanglingxi1993.github.io/dirty_pagetable/dirty_pagetable.html)
- Attacking page tables requires **knowing the physical address** of kernel text/heap
- How about bruteforcing?
  - **No**, I can trigger UAF around **5** times before the kworker dies — not enough

# How About Dirty Page Table Attack?

- Good description: [yanglingxi1993.github.io/dirty\\_pagetable/dirty\\_pagetable.html](https://yanglingxi1993.github.io/dirty_pagetable/dirty_pagetable.html)
- Attacking page tables requires **knowing the physical address** of kernel text/heap
- How about bruteforcing?
  - **No**, I can trigger UAF around **5** times before the kworker dies — not enough
- How about a KASLR infoleak from `msg_msg` out-of-bounds read?
  - **Ok**, let's give it a try!

# Physical Versus Virtual KASLR

- VM run #1

```
gef> ksymaddr-remote
[+] Wait for memory scan
0xffffffff98400000 T _text

gef> v2p 0xffffffff98400000
Virt: 0xffffffff98400000 -> Phys: 0x57400000
```

- VM run #2

```
gef> ksymaddr-remote
[+] Wait for memory scan
0xffffffff81800000 T _text

gef> v2p 0xffffffff81800000
Virt: 0xffffffff81800000 -> Phys: 0x18600000
```

- Virtual address minus physical address:

- VM run #1:  $0xffffffff98400000 - 0x57400000 = 0xffffffff41000000$
- VM run #2:  $0xffffffff81800000 - 0x18600000 = 0xffffffff69200000$

- $0xffffffff41000000 \neq 0xffffffff69200000$

- Sorry, leaking the virtual KASLR offset **doesn't help** against the physical KASLR

# Physical KASLR Versus Virtual KASLR



Physical KASLR



Virtual KASLR

imgflip.com

# Still Needed to Invent Arbitrary Address Writing Primitive

## ① Dirty Page Table Attack?

- Requires page allocator feng-shui to leak the kernel physical address
- No, would be too complicated

## ② Turn UAF write to some kernel object into arbitrary address writing?

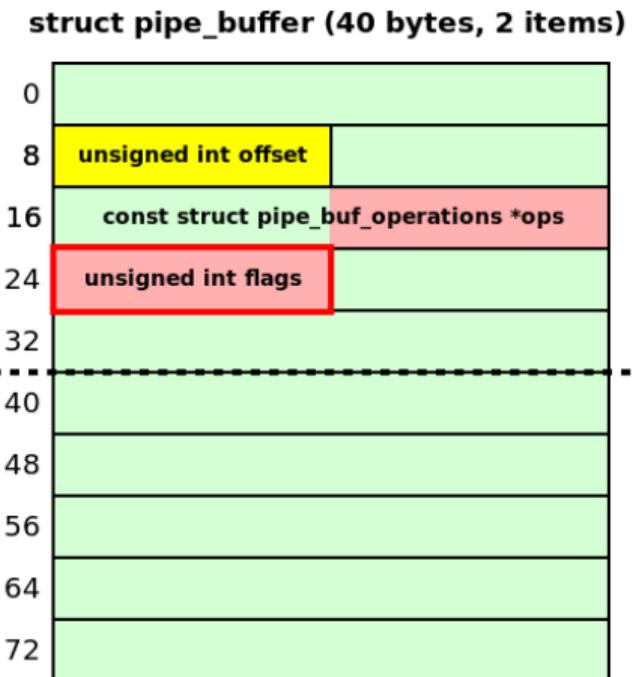

- Not so easy... Exhausting!
- Looked through dozens of different kernel objects
- Read dozens of kernel exploit write-ups
- Tried [Kernel Exploitation Dashboard](#) by Eduardo Vela & KernelCTF team
- Then focused on [pipe\\_buffer](#) kernel object



## Target for UAF Write: struct pipe\_buffer

- We can make `pipe_buffers` of similar size with `virtio_vsock_sock`:
  - Relocate the write end of the pipe
  - `fcntl(pipe_fd[1], F_SETPIPE_SZ, PAGE_SIZE * 2);`
  - The object size becomes: `2 * sizeof(struct pipe_buffer) = 80`
  - Suitable for `kmalloc-96`, like `virtio_vsock_sock`
- Attacker-controlled bytes of `vsock` UAF write change `pipe_buffer.flags`
- It's the original **Dirty Pipe attack** by Max Kellermann [dirtypipe.cm4all.com](http://dirtypipe.cm4all.com)
- Even doesn't need an infoleak
- **One shot, wow, let's try!**

# Target for UAF Write: struct pipe\_buffer



# First of All, Drill!

- Created a Dirty Pipe prototype in **kernel-hack-drill**
- See the code: [kernel-hack-drill/drill\\_uaf\\_write\\_pipe\\_buffer.c](#)
  - Performs cross-cache attack: reclaims `drill_item_t` as `pipe_buffers`
  - Exploits UAF write to `drill_item_t` struct:
    - ★ Controlled bytes at offset **24**
  - Attacker-controlled bytes modify `pipe_buffer.flags`
  - Implements the Dirty Pipe attack
  - LPE in **one shot** without infoleak





<https://www.pngall.com/wp-content/uploads/4/Drill-Machine-PNG-Free-Download.png>

# Not So Fast: CVE-2024-50264 Limitations

- ① Vulnerable `virtio_vsock_sock` client object is allocated together with the server one
- ② Reproducing this race condition is very unstable
- ③ UAF write happens in kworker within few `ps` after `kfree()`
- ④ Null-ptr-deref happens in kworker right after UAF write
- ⑤ If this kernel oops is avoided, another null-ptr-deref happens  
in kworker after `VSOCK_CLOSE_TIMEOUT` (8 sec)
- ⑥ Kworker hangs if `virtio_vsock_sock.tx_lock` is not zero



# Target for UAF Write: struct pipe\_buffer



# Target for UAF Write: struct pipe\_buffer

I can do `splice()` from file to pipe starting from zero `offset` to bypass **Limitation #6!**

**struct virtio\_vsock\_sock (80 bytes)**

0	
8	<code>spinlock_t tx_lock==0</code>
16	
24	<code>u32 peer_buf_alloc</code>
32	
40	
48	
56	
64	
72	



**struct pipe\_buffer (40 bytes, 2 items)**

0	
8	<code>unsigned int offset==0</code>
16	<code>const struct pipe_buf_operations *ops</code>
24	<code>unsigned int flags</code>
32	
40	
48	
56	
64	
72	



# Target for UAF Write: struct pipe\_buffer (No Way)

Oh no, `pipe_buffer.ops` gets corrupted by 4 zero bytes of `peer_fwd_cnt`!

**struct virtio\_vsock\_sock (80 bytes)**

0	
8	<code>spinlock_t tx_lock==0</code>
16	
24	<code>u32 peer_buf_alloc</code>
32	
40	
48	
56	
64	
72	



**struct pipe\_buffer (40 bytes, 2 items)**

0	
8	<code>unsigned int offset==0</code>
16	<code>const struct pipe_buf_operations *ops</code>
24	<code>unsigned int flags</code>
32	
40	
48	
56	
64	
72	

kernel crash



## Target for UAF Write: struct pipe\_buffer (No Way)

- Oh no, `pipe_buffer.ops` gets corrupted by 4 zero bytes of `peer_fwd_cnt`!
  - Changing `peer_fwd_cnt` requires sending data through vsock
  - But successful `vsock connect()` makes the UAF **impossible**
  - **No way** to execute the original **Dirty Pipe attack** 😞

## Target for UAF Write: struct pipe\_buffer (No Way)

- Oh no, `pipe_buffer.ops` gets corrupted by 4 zero bytes of `peer_fwd_cnt`!
  - Changing `peer_fwd_cnt` requires sending data through vsock
  - But successful `vsock connect()` makes the UAF impossible
  - No way to execute the original Dirty Pipe attack 😞
- Suddenly I got a bright idea

# What If?

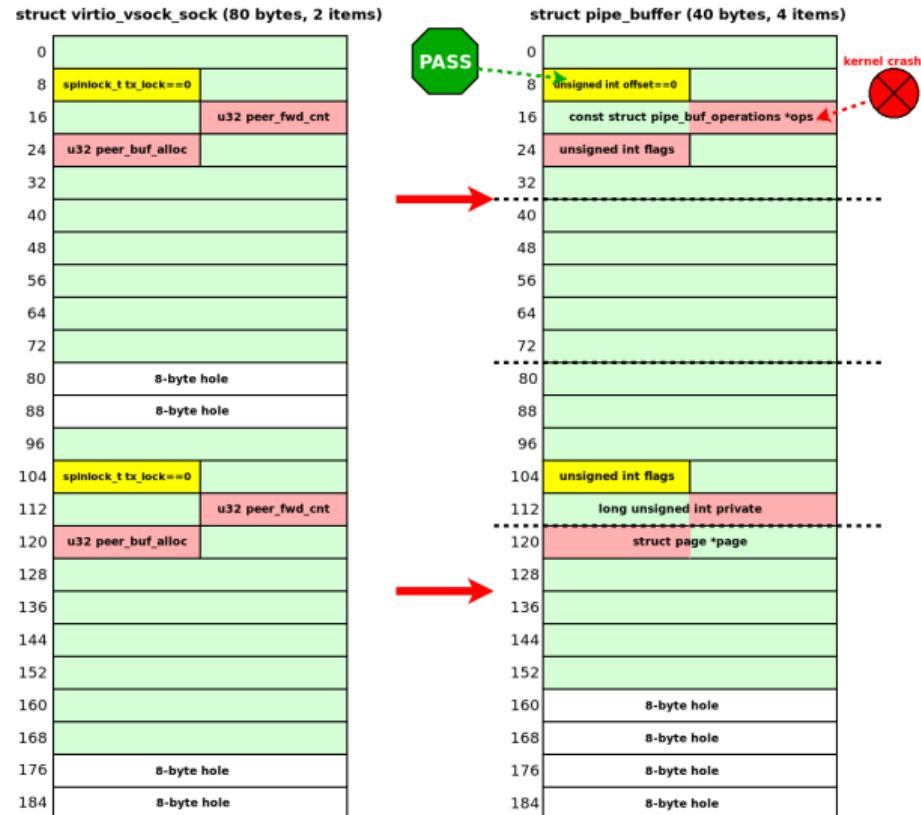


New hope

What if I allocate 4 `pipe_buffers` in `kmalloc-192`?

# Target for UAF Write: Four pipe\_buffers


- Oh no, `pipe_buffer.ops` is corrupted by 4 zero bytes!



# Target for UAF Write: Four pipe\_buffers

- Oh no, `pipe_buffer.ops` is corrupted by 4 zero bytes!
- The kernel crashes if I read from the pipe
- Idea: I discarded the first `pipe_buffer` before UAF
- In that case the bad `pipe_buffer.ops` isn't used!
- How to do it without changing offset:

```
splice(pipe_fds[i][0], NULL,  
      temp_pipe_fd[1], NULL, 1, 0);  
  
read(temp_pipe_fd[0],  
     pipe_data_to_read, 1);
```




# Target for UAF Write: Four pipe\_buffers

- Made `flags` of `pipe_buffer` #3 zero by using `splice()` from file

```
splice(temp_file_fd, &file_offset,  
      pipe_fds[i][1], NULL, 1, 0);
```

- [+] Corrupted `pipe_buffer.page!` YES!
- `kernel-hack-drill` helped to develop it



# Last Revenge From Physical KASLR

- We don't know where the kernel text is inside `vmemmap`
- We can't point `pipe_buffer.page` to kernel code 😞
- Let's shoot to the leaked `struct cred` in the kernel heap
- I can calculate the offset of `struct page` pointing to `cred`:

```
#define STRUCT_PAGE_SZ 64lu
#define PAGE_ADDR_OFFSET(addr) (((addr & 0xfffffffflu) >> 12) * STRUCT_PAGE_SZ)
uaf_val = PAGE_ADDR_OFFSET(cred_addr);
```

- Don't need to know the `vmemmap_base`!
  - [!] I overwrite only 4 lower bytes of `pipe_buffer.page`
- Randomized `vmemmap_base` address has only 2 random bits in lower bytes

## Bruteforce 2 Bits

- In case of **fail** reading from pipe simply returns "**Bad address**"
- In case of **success** reading from pipe gives **struct cred** contents



- Finally, I write zero pipe, overwrite **euid** and **egid**, and **I AM ROOT**

## Demo Time



# Conclusion

- Bug collision is painful
- But finishing the research anyway is rewarding
- Try my open source project  
[github.com/a13xp0p0v/kernel-hack-drill](https://github.com/a13xp0p0v/kernel-hack-drill)
- **kernel-hack-drill** is a useful testing environment  
for Linux kernel security researchers
- Contributors are always welcome!



# Thanks 감사합니다

## Enjoy the conference!

Contacts:



 a13xp0p0v@tuta.io

Blog: [a13xp0p0v.github.io](https://a13xp0p0v.github.io)



Channel: [t.me/linkersec](https://t.me/linkersec)

